

Mechanized Semantics *for the Clight* Subset *of the C* Language

J Autom Reasoning (2009) 43:263–288
DOI 10.1007/s10817-009-9148-3

**Mechanized Semantics for the Clight Subset
of the C Language**

Sandrine Blazy · Xavier Leroy

hello_world.c



clang



a.out

hello_world.c

This “does” X


clang

a.out

Does this “do” X?


hello_world.c

This “does” X



clang

Does clang preserve X?



a.out

Does this “do” X?

$X \approx$ “observable behaviors” of the program



I/O Events
(e.g. system calls)



Result
(e.g. non-termination)

hello_world.c

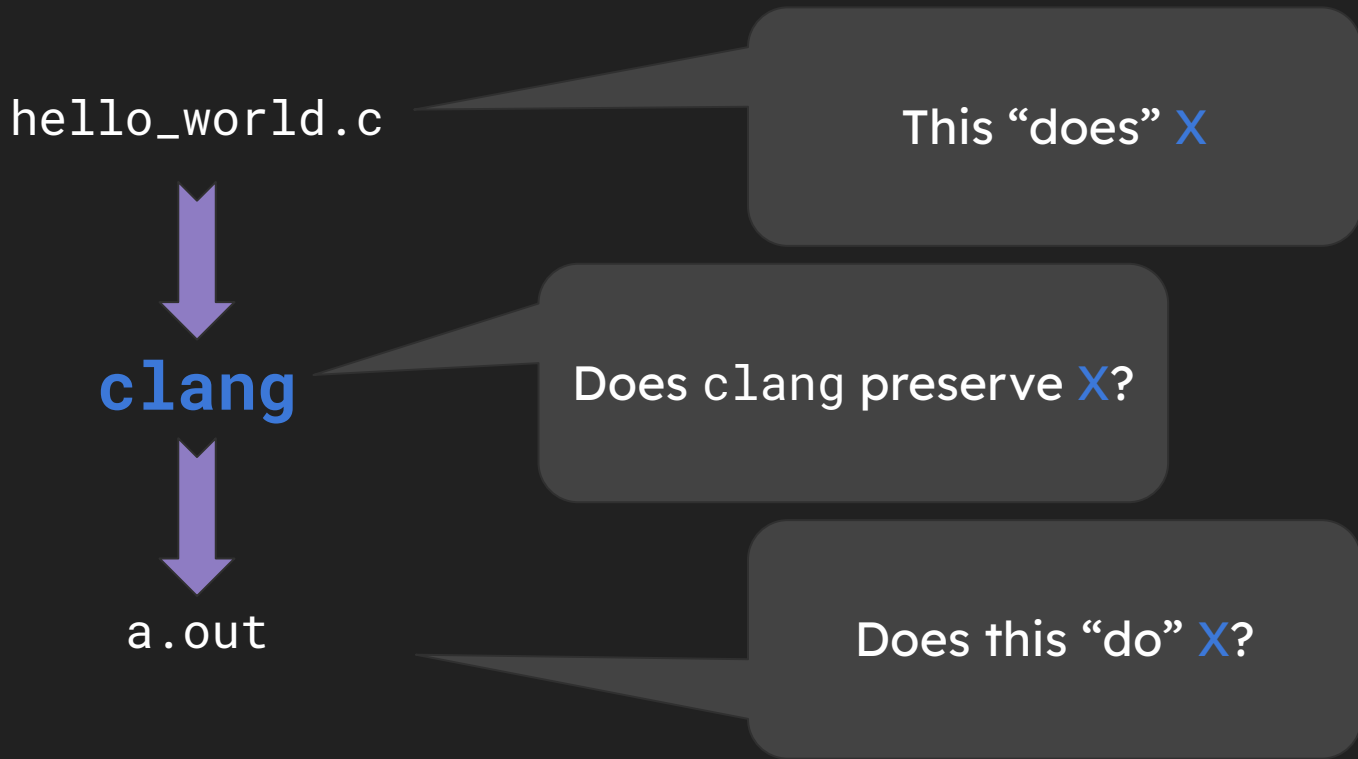
This “does” X

clang

Does clang preserve X?

a.out

Does this “do” X?



hello_world.c

This program's
observable
behavior is **X**

clang -O1 (with bugs)

a.out

This program's
observable
behavior is **Y**

Exit Status of This Program?

```
int x = 0;

int test(int* restrict ptr) {
    *ptr = 1;
    if (ptr == &x)
        *ptr = 2;
    return *ptr;
}

int main() { return test(&x); }
```


Exit Status of This Program?

```
int x = 0;

int test(int* restrict ptr) {
    *ptr = 1;

    if (ptr == &x)
        *ptr = 2;

    return *ptr;
}

int main() { return test(&x); }
```

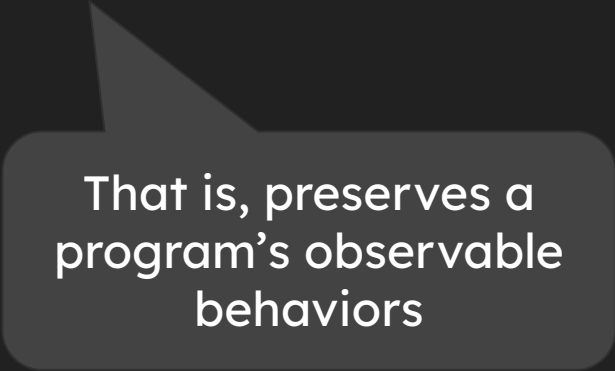
```
> clang -O1 err3.c && ./a.out
Clang 17.0.6
test(&x) = 1
```

```
> gcc -O3 err3.c && ./a.out
13.2.1 20240417
test(&x) = 2
```

```
> ccomp -O3 err3.c && ./a.out
314
test(&x) = 2
```

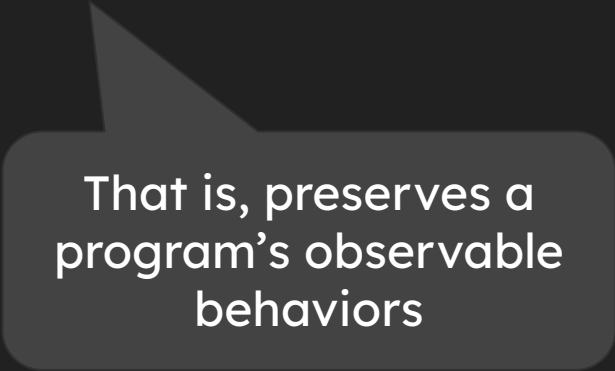
miscompilation

How do we know if our compiler **preserves semantics**?



That is, preserves a
program's observable
behaviors

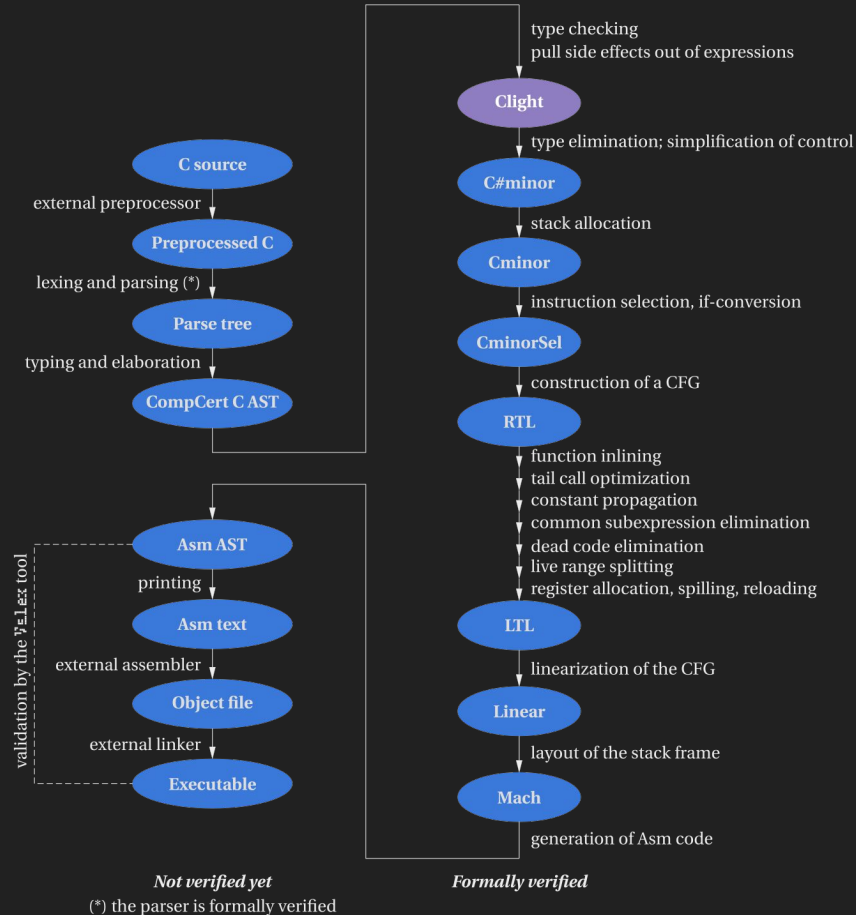
How do we know if our compiler **preserves semantics**?



That is, preserves a
program's observable
behaviors

Formal verification.

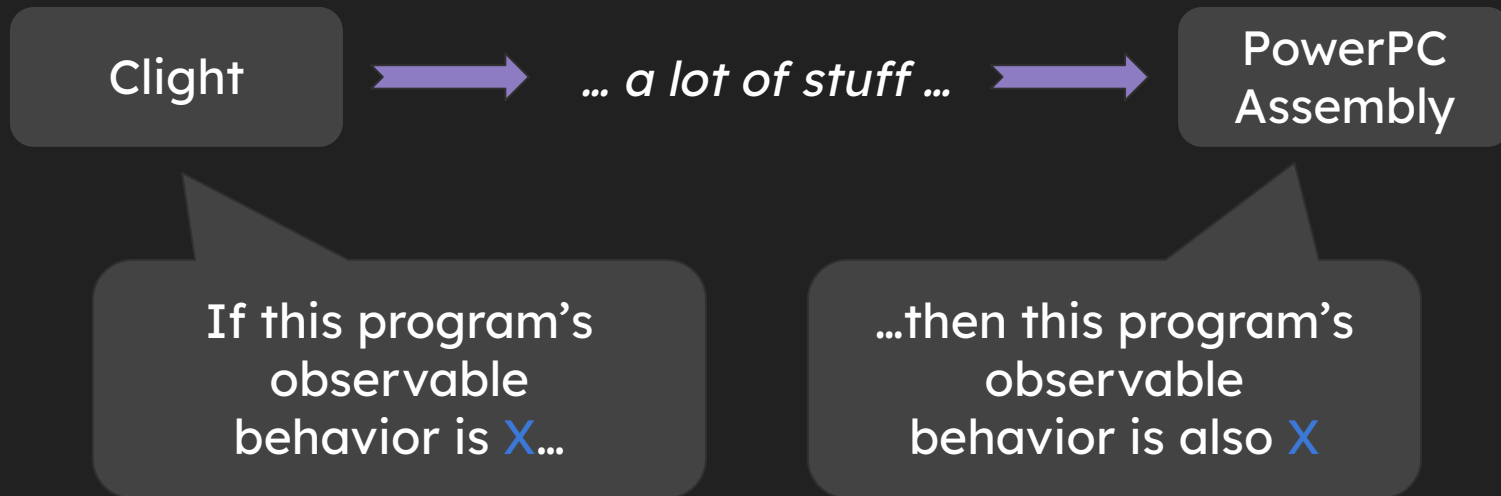
CompCert \approx 90% verified compiler



CompCert \approx 90% verified compiler



CompCert \approx 90% verified compiler

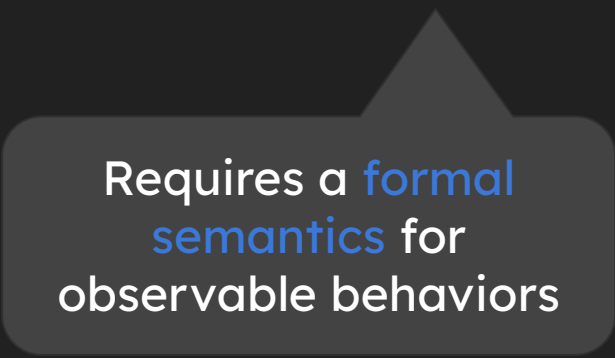


How do we know CompCert preserves semantics?

By proving that the observable behaviors
before and after compilation are the same.

How do we know CompCert preserves semantics?

By proving that the observable behaviors
before and after compilation are the same.



Requires a **formal semantics** for
observable behaviors

How do we know CompCert preserves semantics?

By proving that the observable behaviors
before and after compilation are the same.

Requires a **formal semantics** for
observable behaviors


Which itself is going to
require a **formal semantics**
for the source language

Mechanized Semantics *of the* C Language


J Autom Reasoning (2009) 43:263–288
DOI 10.1007/s10817-009-9148-3

**Mechanized Semantics
of the C Language**

Sandrine Blazy · Xavier Leroy



Undefined and unspecified behaviors...



Undefined and unspecified behaviors...

...don't play nice with formal semantics

Mechanized Semantics *for the Clight* Subset *of the C* Language

J Autom Reasoning (2009) 43:263–288
DOI 10.1007/s10817-009-9148-3

**Mechanized Semantics for the Clight Subset
of the C Language**

Sandrine Blazy · Xavier Leroy

This paper defines a formal semantics
for the Clight language

Clight is a subset of C that
seeks to reduce C's ambiguity

J Autom Reasoning (2009) 43:263–288
DOI 10.1007/s10817-009-9148-3

Mechanized Semantics for the Clight Subset of the C Language

Sandrine Blazy · Xavier Leroy

How does Clight formalize
observable behaviors?

I/O Events

Traces of external function calls are tracked

Clight statements may update the trace

$$G \vdash F(v_{\text{args}}), M \xrightarrow{\text{T}} \infty$$

Results

Programs may **terminate** with a final value or **diverge**

$$\vdash P \Rightarrow \text{terminates}(t, n)$$

$$\vdash P \Rightarrow \text{diverges}(T)$$

How does Clight handle
ambiguous C features?

How do we evaluate the following code snippet?

a +++ b

(a++) + b

a + (++b)

a +++ b

$(a++) + b$

$a + (++b)$

These are expressions

a +++ b

(a++) + b

a + (++b)

These are expressions

But they also have a side effect...

In C, expressions can have side effects:

`a * b++`

...and the evaluation order of these expressions is sometimes ambiguous:

`a +++ b`

In Clight, expressions **cannot** have side effects

~~$a * b + c$~~

...which gets rid of ambiguous evaluation orders:

$a + b$

In Clight, expressions **cannot** have side effects

- No function calls
- No assignments (++ , += , etc.)

These features must appear in Clight **statements**

In Clight, expressions **cannot** have side effects

- Simplifies semantics for expressions
- Allows expressions to be evaluated in any order

Evaluation order partially
unspecified in C

How do we evaluate the following code snippet?

`a +++ b`

`(a++) + b`

`a + (++b)`

How do we evaluate the following code snippet?



Trick question. This program is incorrect.

How do we evaluate the following code snippet?

```
*a + fun(a)
```

How do we evaluate the following code snippet?



```
*a + sin(a)
```

Trick question. This program is incorrect.

Other examples...

What is the size (bytes) of the following variable?

```
long a = 0;
```


What is the size (bytes) of the following variable?

```
long a = 0;
```

4? 8? It depends...

Sizes of types in C can
be ambiguous

Clight uses exact, specified sizes for its types...

What is the size (bytes) of the following variable?

```
I32 a = 0;
```



4.

What about other
interesting C features?

Pointers become **locations**

Locations come in the form (b, δ)

Block within
memory

Offset
within block

Pointers become **locations**

Locations come in the form (b, δ)

Pointer arithmetic simply adds to δ in (b, δ)

Pointers become **locations**

Locations come in the form (b, δ)

Pointer arithmetic simply adds to δ in (b, δ)

Functions and arrays become locations when passed as arguments

Variables can be in:

- Global scope
- Function scope

Clight does not have **block** scope

Variables can be in:

- Global scope
- Function scope

Clight does not have **block** scope

Simplifies the retrieval of a variable's address with **&**

How were Clight's
semantics **verified**?

Manual review by C compiler experts

Manual review by C compiler experts

Perform tests using an interpreter (*in progress*)

Manual review by C compiler experts

Perform tests using an interpreter (*in progress*)

Define multiple semantics for Clight
and prove their equivalence (*in progress*)

We can use other verified parts of the CompCert compiler

Clight  X

If we verify properties about X...

...then we implicitly verify properties about Clight

Thank you for listening!